



PIGEON LINK

Our Infrastructure, Your AI Agents

But Deploying AI agents is costly and complicated



Scaling Costs Are Unsustainable

AI infrastructure spending will hit \$421B by 2033, making traditional scaling too expensive.



AI Agents Struggle to Communicate

60% of AI projects struggle with integration due to lack of standardized protocols for deployment.



Multi-Agent AI Is Hard to Coordinate

48% of energy and 42% of telecom firms report failures when trying to scale AI due to poor coordination.



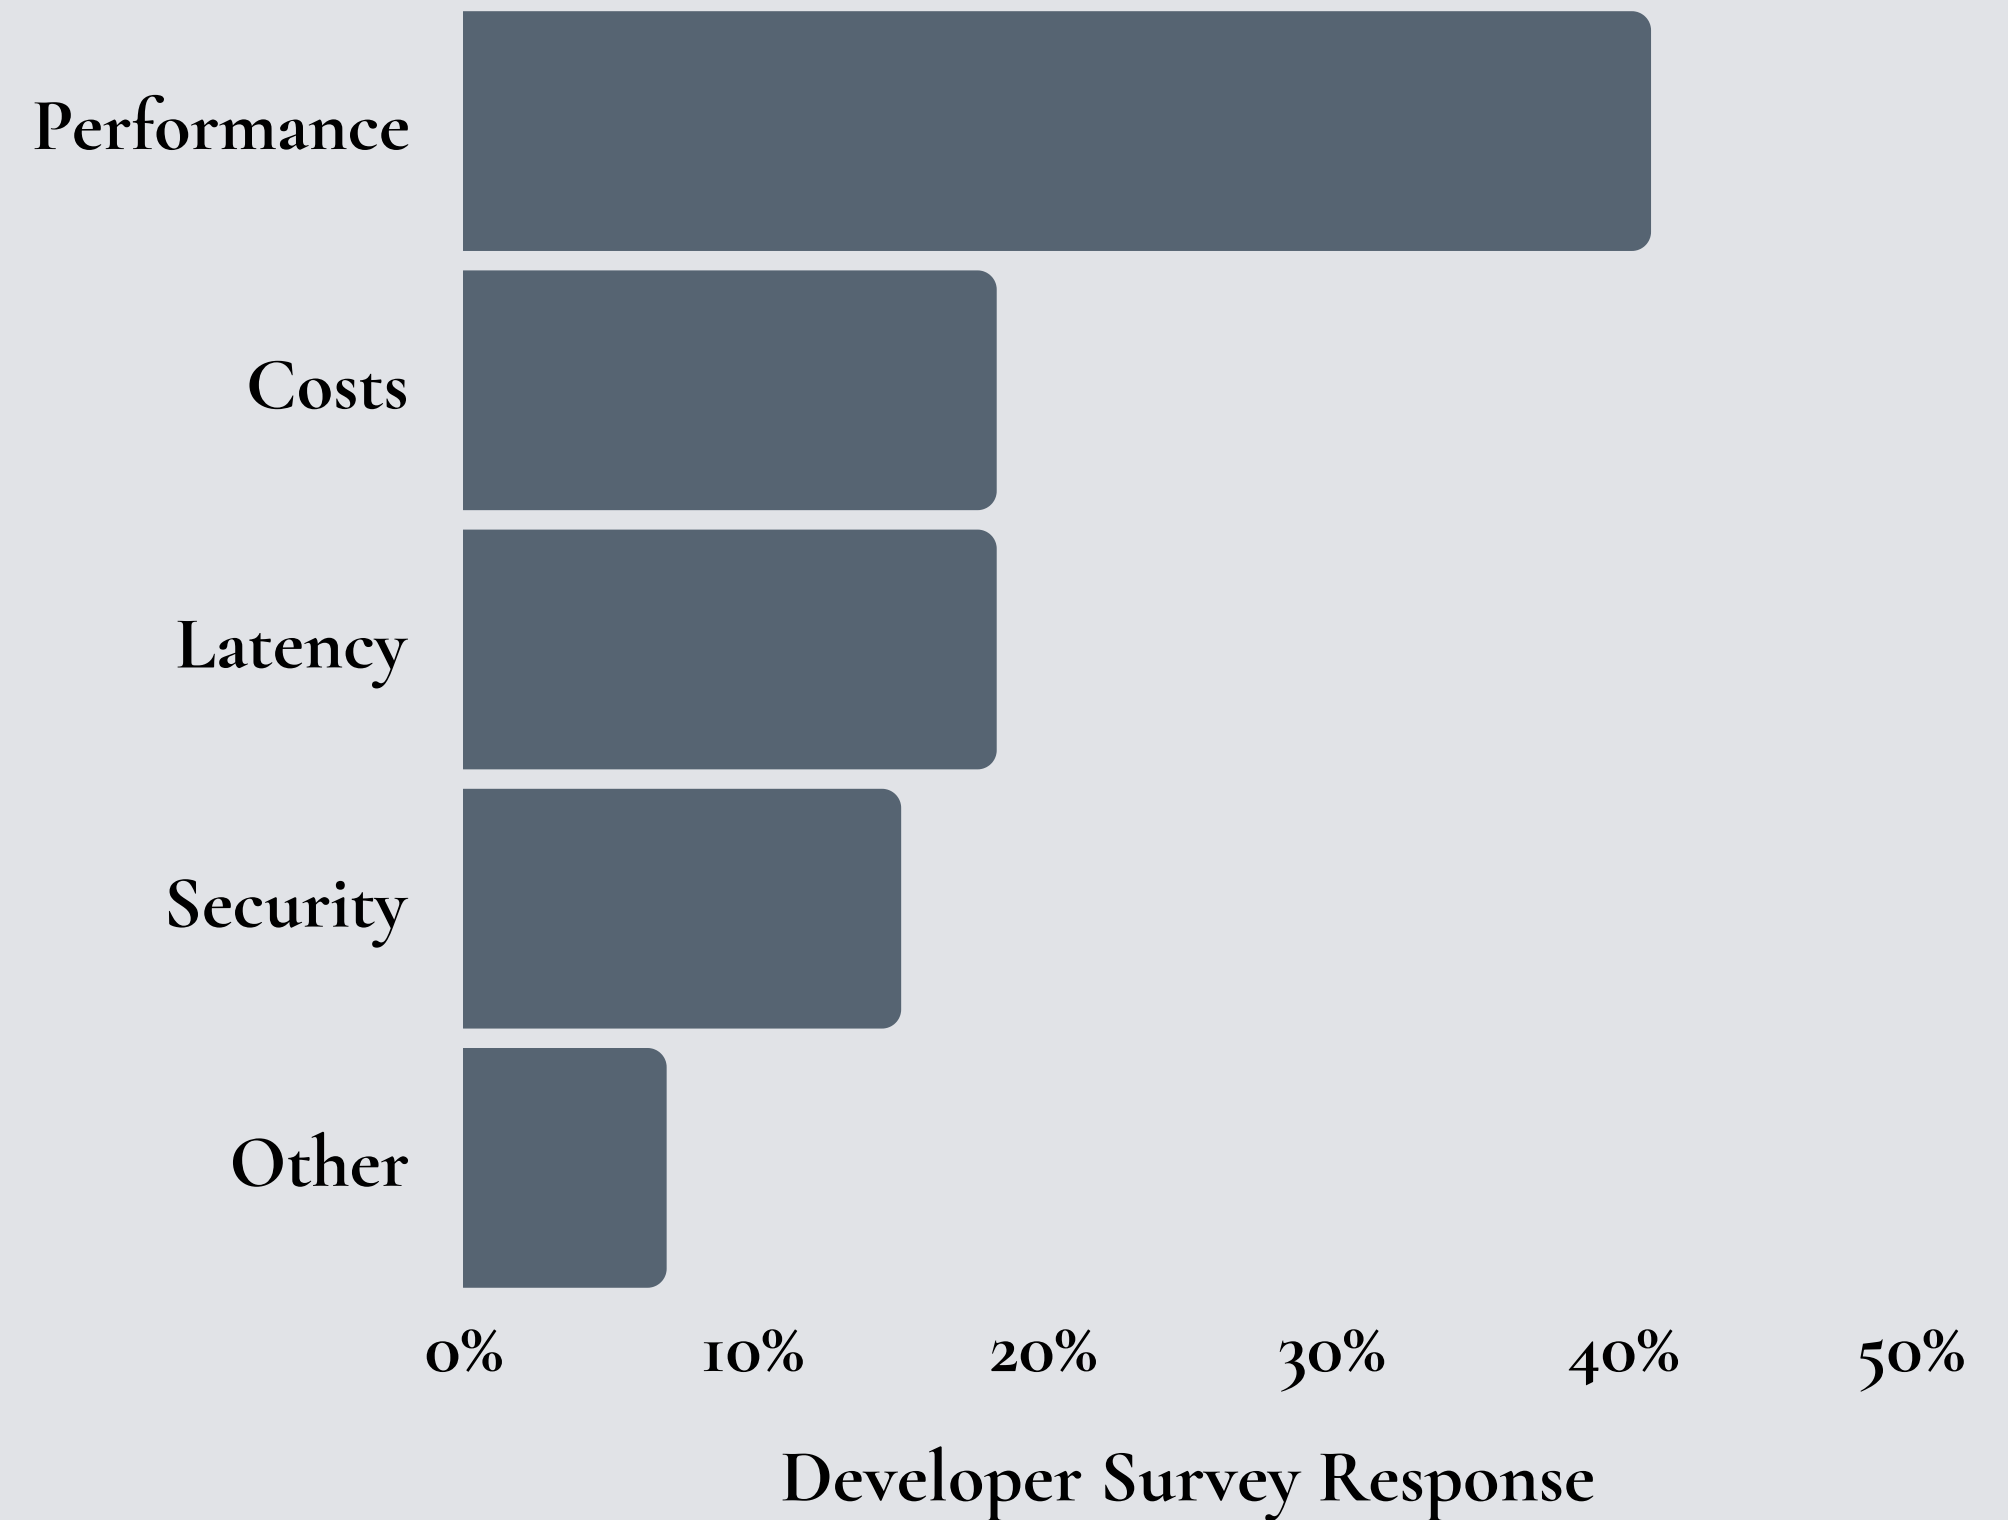
Debugging AI Is a Nightmare

40% of AI projects fail because tracing requests across agents is nearly impossible

“It feels like many of these issues come from the ecosystem moving too fast. Still, simplicity in DX like deploying on DO/Vercel just feels missing.”

- a developer deploying AI Agents

Biggest limitation of putting more agents in production



Pigeon fixes AI Agent Deployment

Distributed Inference

ensures smooth workflows across servers

Optimized Hardware

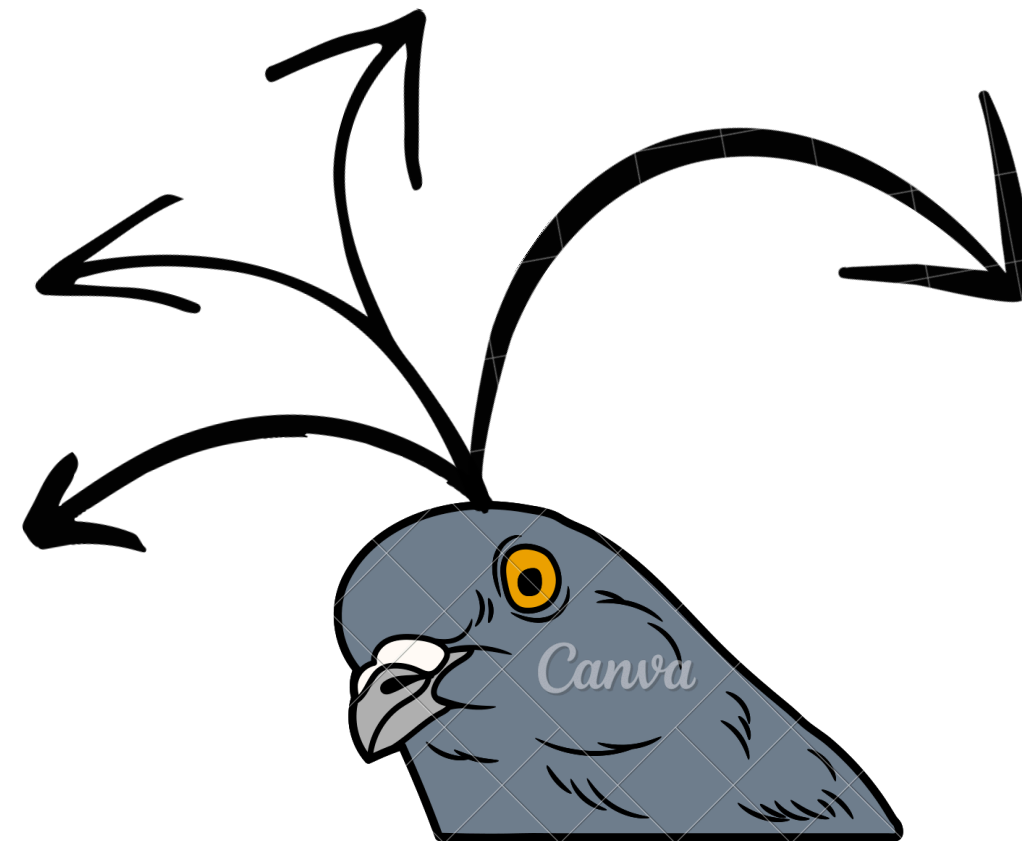
and OS/kernel support
for better performance

Ease of programming

expressive meta-language
to program agents

Communication Protocol between Agents

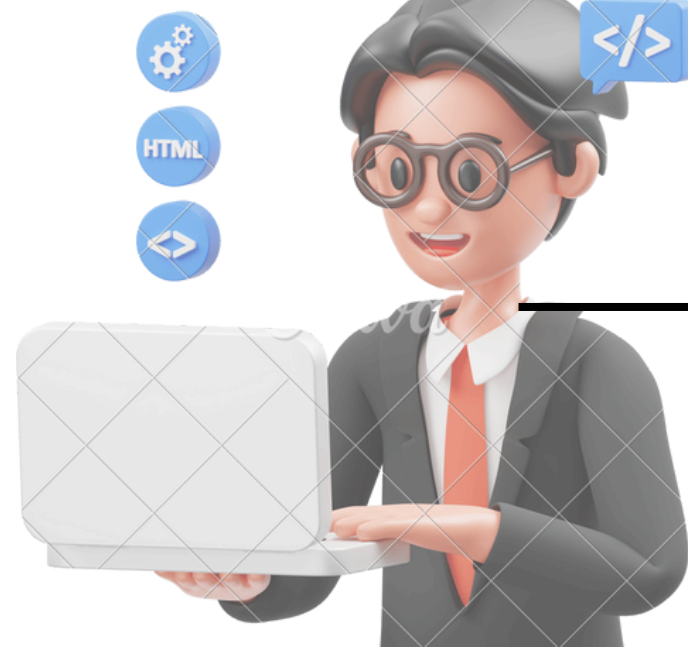
ensures reliable interactions
within the agentic network



Pigeon reduces scaling costs and lets you focus on building agents, not infrastructure.

How Pigeon supports Industrial Automation

Factory IT team
creates an app layer
to automate assembly
line operations.



Factory IT team deploys
AI agents for **quality
inspection,**
maintenance, and
inventory on Pigeon



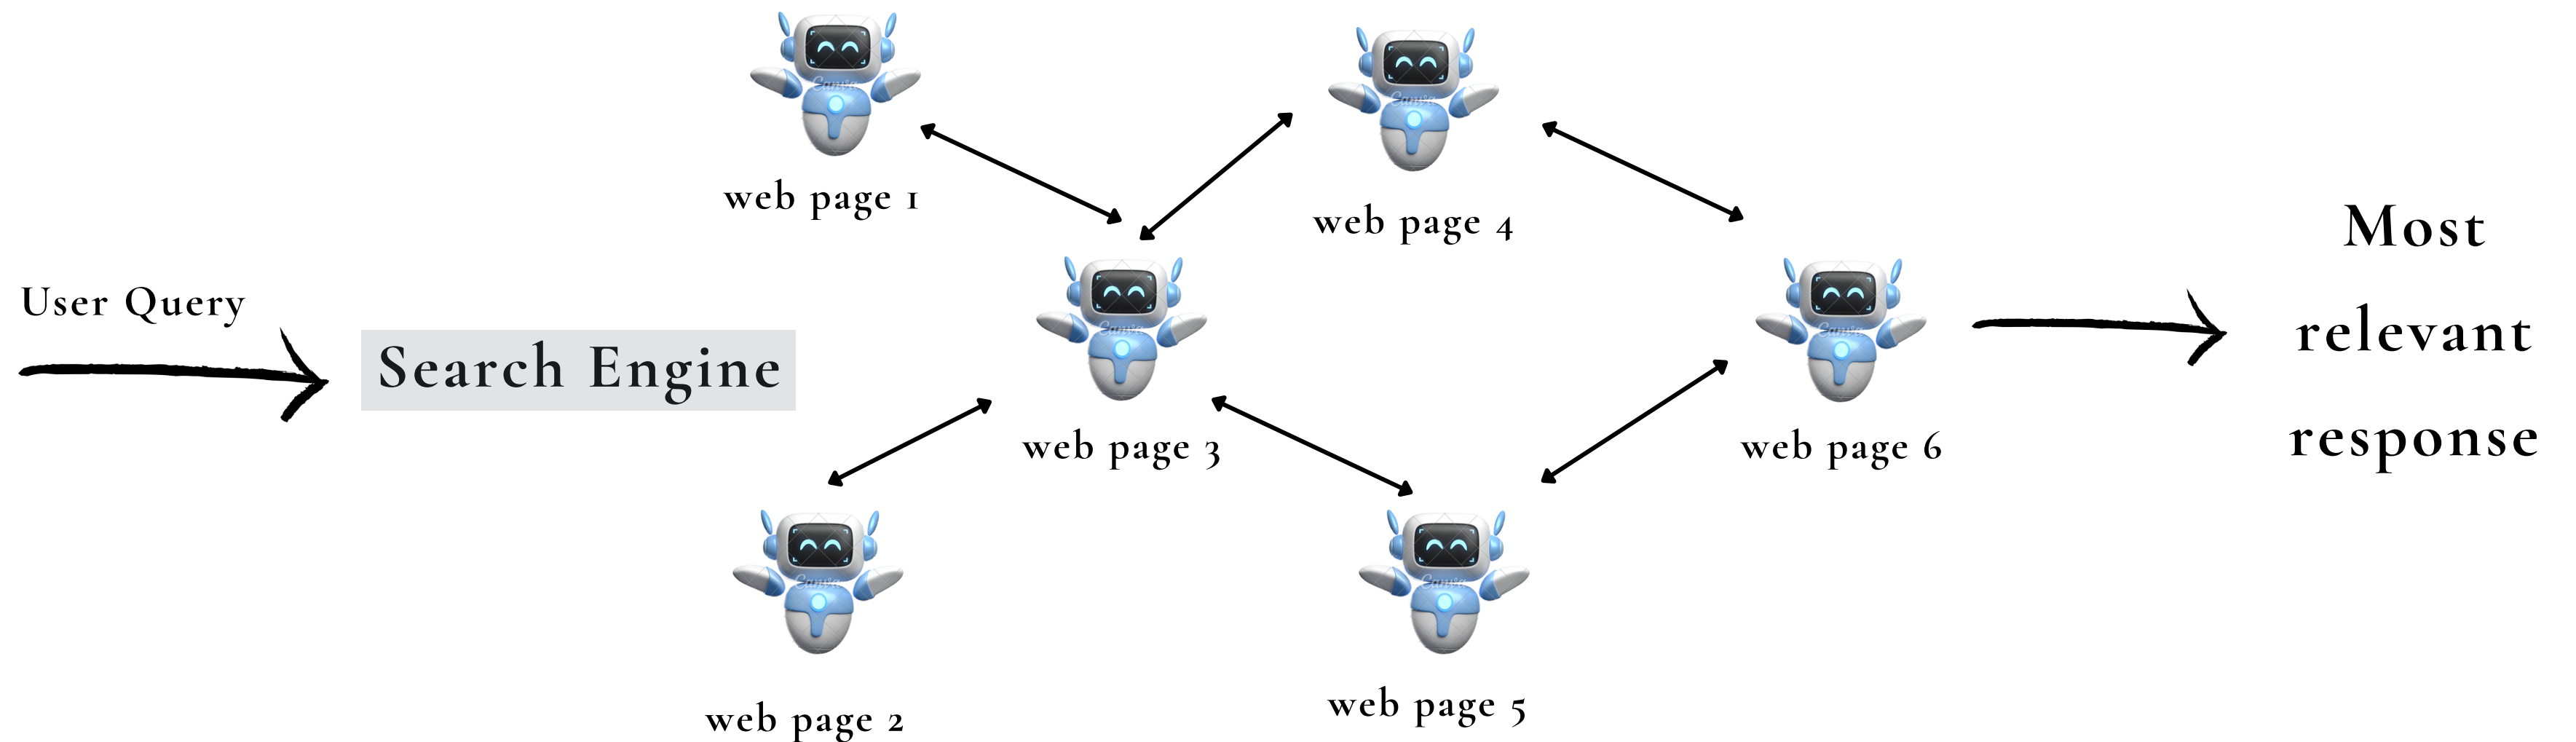
Application Layer

Scalable infrastructure that grows
with the factory floor,

A communication network
between the Agents to work
together

Low-latency execution with
automatic agent scaling, handling
workload routing without manual
intervention

The future of Agentic Internet



Why Now?

93%

of IT leaders plan
autonomous agents; nearly
half already have.

73%

increase in AI models
predicted over the next
three years

86%

of Enterprises Require Tech
Stack Upgrades to Properly
Deploy AI Agents

Usage-Based Pricing

Flexible Pricing

Pay-as-you-go: Customers pay based on actual agent usage

Scales automatically as customer needs grow.

Team

Vidhi Kothari

CEO

Formerly at Morgan Stanley, at Newmark, and Quadreal, specialized in investing. Most recently, focused on go-to-market strategy at Copia Automation (Series A)

Rajpreet Singh

CTO

As a Research Intern in the Autonomous Systems Group at Bell Labs, Munich, and a graduate student at TU Munich specializing in High-Performance Computing (HPC), Rajpreet developed the idea while deploying workflows for an internal project.